

Mr Yizhou Wang

Contact No.: +86 18215557703 | Email: wangyizhou317@gmail.com

HKUSTGZ NLP Group | Personal Page: lzhou-wang.github.io

Education

09/2019-07/2023

University of Leeds

- Bachelor of Engineering in Electronic and Electrical Engineering
- Graduated as First Class I with Honour

Southwest Jiaotong University (西南交通大学)

- Bachelor of Engineering in Electronic and Electrical Engineering
- GPA: 3.55/4
- Professional Skills: MATLAB, Python <PyTorch, Transformers>

08/2023-Present

The Hong Kong University of Science and Technology (Guangzhou)

- Master of Philosophy (M.Phil.) Supervised by Prof. [Song JIE](#) and Prof. [Xuming HU](#)
- Postgraduate Scholarship of ¥ 240,000
- Research on Visual Language Model Hallucination Mitigation and Role-LLM Reasoning
- Current GPA: 3.56/4

Internship

08/2025-Present

[Shanghai AI Lab](#)

- Position at Knowledge Frontier Centre
- Research on large scale multi-model information retrieval
- Research on multi-encoder MLLMs
- Research on RL for MLLM reasoning

Project

10/2025-Present

[Ongoing] Mitigating Over-thinking in Language Model Reasoning via Token-stage Reward

- Reducing thinking length for language model reasoning
- Modified GRPO RL training framework for efficient reasoning
- Project Repository available soon

05/2025-Present

[Ongoing] Edge Device STT-LLM-TTS Demo

- Deploy fine-tuned LLM with STT and TTS on RK35**-based system
- Extremely low latency for edge on-device chat experience
- Project Repository available soon

02/2025-Present

[Ongoing] YouShouChat: Role-Play LLM with Unique Characteristics and Memories

- One-step workflow for data generation and sampling
- Customized training strategies, quantization and RAG framework
- Project Repository available at <https://github.com/lzhou-Wang/YouShouChat>
- Currently ~~open for test~~ offline

08/2025-Present

[Under Review] Investigating Redundancy in Multimodal Large Language Models with Multiple Vision Encoders

- Published as first author

- Research on encoder-level efficiency for MLLMs with multiple vision encoders
- Experimental results show large information gap between encoders
- MLLMs do not behave well with more than 2 vision encoders
- Pre-Print Paper available at <https://arxiv.org/pdf/2507.03262>

05/2025-Present

[Under Review] **CondenseVLM: High-Fidelity Visual Token Condensation for Large Vision-Language Models**

- Developed token merging for MLLM efficient reasoning
- SOTA performance in 9 popular benchmarks
- Project Repository available soon
- Pre-Print Paper available soon

02/2025-05/2025

[NeurIPS 2025] **Don't Just Chase "Highlighted Tokens" in MLLMs: Revisiting Visual Holistic Context Retention**

- Developed holistic token pruning for MLLMs
- SOTA performance in 9 popular image-QA benchmarks at different pruning ratios
- SOTA performance in 2 popular video-QA benchmarks
- More efficient and reasonable pruning strategies upon visualization
- Project Repository available at: <https://github.com/obananas/HoloV>
- Pre-Print Paper available at: <https://arxiv.org/pdf/2510.02912>

07/2024-11/2024

[ICML 2025] **Memory-Space Visual Retracing for Hallucination Mitigation in Multimodal Large Language Models**

- Published as co-first author
- Developed training-free inferencing structure for mitigating MLLM hallucination
- SOTA performance in 7 popular benchmarks with 3 MLLMs tested
- Extremely high efficiency in inferencing token generation speed and memory occupation
- Popular GitHub Repository with 100+ stars!
- Project Repository available at <https://github.com/1zhou-Wang/MemVR>
- Pre-Print Paper available at <https://arxiv.org/pdf/2410.03577>

12/2023-Present

Multi-Functional Mobile Platform

- Awarded with Excellent RBM Project (Top 8/70)
- Developed Mobile robot platform with gripper and dexterous hand
- ICRA 2024 AXS Sim2Real Challenge 3rd in Simulation stage and 2nd in onsite competition
- Project Repository available at: <https://github.com/1zhou-Wang/Multifunctional-Mobile-Platform>

Service:

10/2025-Present

ICLR 2025 Reviewer